


Fugusashi Benchmark Results (v1.3.0)



Metrics: 30 held-out prompts, 3 model classes (gpt-oss-120b, hermes-3-405b, lfm-2.5-1.2b)

Router latency: 83ms (ModernBERT) | Training: 224 samples, 137s | Federated: 3 clients with DP ($\epsilon=1.8$)

Fugu (Sakana AI) baseline: proprietary 7B coordinator LLM, trained on undisclosed data