


Fugusashi Routing Accuracy (v1.3.0)



Metrics: 30 held-out prompts (6 categories × 5) | 3 model classes (gpt-oss-120b, hermes-3-405b, lfm-2.5-1.2b)

Router latency: 22ms median (M3 Pro) | Training: 224 samples, 163s | Fisher's exact test $p = 1.4 \times 10^{-3}$

*Preliminary simulated federated evaluation on 20 hand-curated prompts — not a published benchmark